
Linux Kernel Basecamp

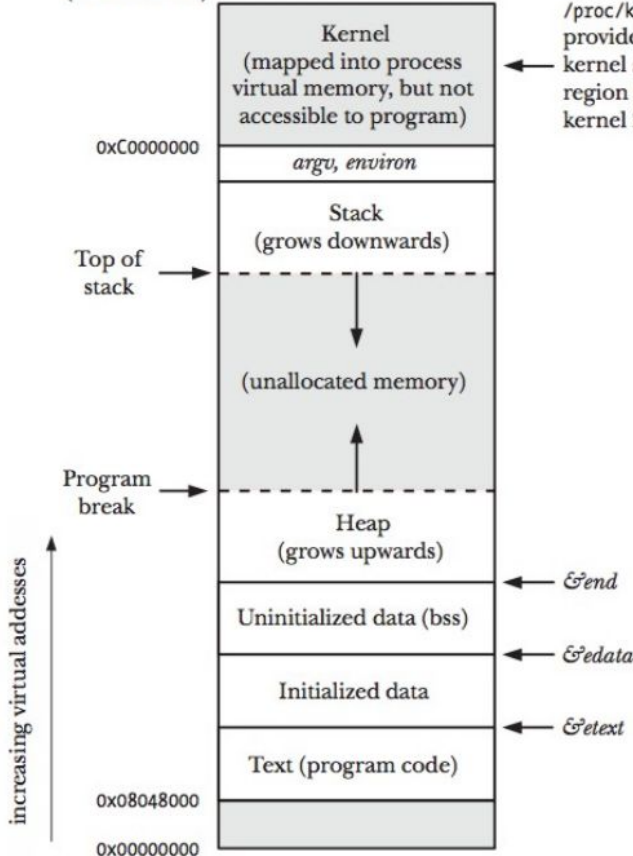
— Memory Management —
Apr 2023

Agenda

- Memory layout
 - Dynamic memory allocation
 - SLAB allocators
 - Mempools
-
-

Memory Layout

Virtual memory address
(hexadecimal)




Memory page frame size on x86-32 is 4096 bytes

`$getconf PAGE_SIZE`

Userspace functions

`malloc()` / `calloc()` / `realloc()` / `free()`
`brk()` / `sbrk()`
`mmap()`
`posix_memalign()`

Overcommitting and OOM



Virtual memory

Userspace memory allocation

Internal and external fragmentation are the common issues

Glibc allocates memory via

- ✓ buddy memory allocation scheme
- ✓ anonymous memory mapping

`MMAP_THRESHOLD` is 128 kB (`malopt()` to change)

Memory Management

```
#include <include/linux/gfp.h> or/and
```

```
#include <include/linux/types.h>
```

```
void *kmalloc(size_t size, gfp_t flags);
```

gfp_t: is a bitmask. The most popular predefined values are:

GFP_ATOMIC, GFP_KERNEL, GFP_USER



```
void kfree(const void *);
```

```
#include/linux/vmalloc.h
```

```
void *vmalloc(unsigned long size);
```

```
void vfree(const void *addr);
```

Vmalloc vs kalloc



Low level memory management

```
#include/linux/gfp.h
```

```
struct_page *alloc_pages(gfp_t gfp_mask, unsigned int order);
```

```
struct page * alloc_page(gfp_t gfp_mask);
```

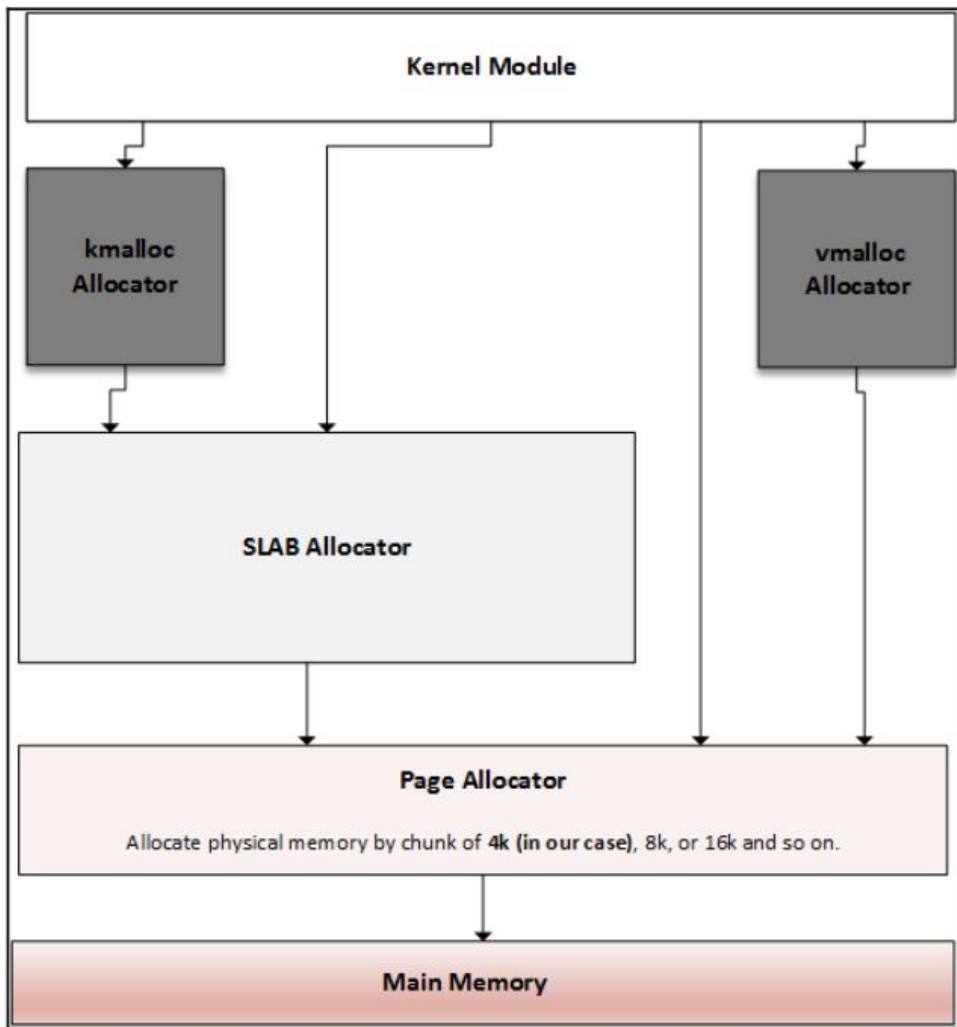
```
void *page_address( struct_page * page );
```

```
unsigned long get_zeroed_page(unsigned int gfp_mask);
```

```
-----
```

```
void free_pages(unsigned long addr, unsigned int order);
```

```
void free_page(unsigned long addr);
```

Kernel memory allocation

MemPool

```
#include/linux/mempool.h

typedef void *(mempool_alloc_t)(int gfp_mask, void *pool_data);

typedef void (mempool_free_t)(void *element, void *pool_data);

mempool_t *mempool_create(int min_nr, mempool_alloc_t *alloc_fn, mempool_free_t *free_fn, void
pool_data);

void *mempool_alloc(mempool_t *pool, gfp_t gfp_mask);

void mempool_free(void *element, mempool_t *pool);

mempool_alloc_slab()

mempool_kmalloc()

mempool_alloc_pages()
```

Interface SLAB

File system point: /proc/slabinfo

include/linux/slab.h

```
struct kmem_cache *kmem_cache_create(const char *name, size_t size, size_t align, unsigned long flags, void (*ctor)(void *));
```

```
void kmem_cache_destroy(struct kmem_cache *s);
```

```
void *kmem_cache_alloc(struct kmem_cache *, gfp_t flags);
```

```
void kmem_cache_free(struct kmem_cache *, void *);
```

- SLAB
- SLUB
- SLOB

```
cat /boot/<path_to_config> | grep "CONFIG_SL[AOU]B"
```